

PROFESSIONAL SUMMARY

Databricks Certified Data Engineer -- Professional and Associate. Senior Data Engineer with **4+ years** of experience building enterprise-scale data platforms for **Fortune 500 and FTSE 100 multinationals** on **Azure Databricks, PySpark, Snowflake** and **AWS**. Currently architecting **medallion (bronze/silver/gold)** pipelines with **Unity Catalog** governance and **Azure Event Hub** streaming ingestion for an **FTSE 100 consumer-health multinational** -- where a data observability and lineage framework I built **reduced daily pipeline failures from ~20% to 1–2%**. Previously delivered a full **SAS-to-Snowflake platform migration** for one of the **world's largest hotel groups**, and introduced **data quality, observability and CI/CD** practices for a **NASDAQ-listed US restaurant group**. Owns architecture decisions end to end, reviews code, mentors junior engineers and runs client requirement conversations directly. Advanced practitioner of **AI-assisted engineering** (Claude Code -- multi-agent workflows), delivering roughly **50% faster** while keeping correctness, performance and security review in the loop. Domain delivery across **consumer goods (CPG), hospitality and QSR** for globally distributed, multi-market businesses.

CORE SKILLS

Databricks & Big Data: Azure Databricks, Apache Spark, PySpark, Spark SQL, Delta Lake, Medallion Architecture (bronze/silver/gold), Unity Catalog (governance & lineage), Databricks Workflows, Databricks AI/BI Genie, Databricks AI functions, Spark performance tuning

Azure: Azure Databricks, Azure Data Factory (ADF), Azure Event Hub, ADLS / Storage Accounts

Data Warehousing: Snowflake (platform migration, stored procedures, performance tuning), Dimensional Modelling (Star/Snowflake schema), SCD Type 1/2

AWS: AWS Glue, Amazon S3

Orchestration & Transformation: Databricks Workflows, Apache Airflow, dbt, ETL/ELT pipeline design

Languages & Libraries: Python, **PySpark**, SQL (advanced -- window functions, query optimisation), Pandas

Practices: Data Governance, Data Quality & Observability, CI/CD (Git-based), Performance & Cost Optimisation, Pipeline Reliability & Root-Cause Analysis, Cloud Migration, Streaming Ingestion, Client Stakeholder Management, Mentoring & Code Review

AI-Assisted Engineering: Claude Code (advanced -- multi-agent workflows, parallel agent orchestration, forked sessions), AI-assisted development & refactoring, review of generated code for correctness, performance and security

Working knowledge: Azure Synapse Analytics, Delta Live Tables (DLT), Databricks Auto Loader

CERTIFICATIONS

Databricks Certified Data Engineer Professional -- Databricks · Issued August 2026
Credential ID 192414204

Databricks Certified Data Engineer Associate -- Databricks · Issued August 2026
Credential ID 192183984 · credentials.databricks.com/fd6b93bc-09f0-437b-8d9d-171e34168eb3

A Crash Course in PySpark -- Udemy
udemy.com/certificate/UC-c89c0ca4-bd7b-478f-b1ab-92e86cff527a

PROFESSIONAL EXPERIENCE

Data Engineer -- Quantzig

FTSE 100 · CONSUMER HEALTH

Hyderabad, India · Oct 2025 – Present

- Cut daily pipeline failure rate from ~20% to 1–2%** for an **FTSE 100 multinational consumer health and hygiene group** selling across nearly 200 markets, by diagnosing recurring root causes and **isolating failure-prone jobs** so a

single failure no longer cascades into dependent workloads.

- Designed and built a **data observability framework** giving end-to-end visibility into pipeline health -- including **edge-case failures that previously went undetected**.
- Built **data lineage linking Power BI reports to upstream pipeline failures**, making the business impact of any failure immediately traceable and cutting triage and stakeholder-communication time.
- Applied **Databricks AI functions to automate root-cause analysis**, surfacing probable failure causes automatically instead of relying on manual log inspection.
- Design and build production pipelines on **Azure Databricks** implementing **medallion architecture** (bronze/silver/gold), curating raw ingested data into analytics-ready gold tables.
- Integrated **five source systems** to power the observability and lineage layer -- **Power BI REST API, Azure Analysis Services (AAS) cube API, Databricks API, Azure Data Factory via Azure Event Hub**, and **SharePoint** -- unifying pipeline, model and report telemetry into a single view.
- Implemented **batch and streaming ingestion** from **Azure Event Hub** and **ADLS storage accounts**, orchestrated through **Databricks Workflows**.
- Deployed **Unity Catalog** for centralised data governance -- access control, data lineage and auditability across the platform -- and enabled business self-service analytics through **Databricks AI/BI Genie**.
- Apply **AI-assisted engineering workflows** (Claude Code -- multi-agent orchestration, parallel agents, forked sessions) across pipeline development, refactoring, debugging and test coverage, cutting delivery time by **~50%** while reviewing every generated change for correctness, performance and security.
- **Own end-to-end architecture decisions** for the data platform; **review code** and set engineering standards; **mentor junior engineers**; and run **client requirement and stakeholder conversations** directly.

Data Engineer -- Sigmoid US QSR CHAIN

Bengaluru, India · Feb 2025 – Sep 2025

- Built and maintained ETL pipelines for a **NASDAQ-listed US quick-service restaurant group** operating over 2,000 locations nationwide, using **AWS Glue, Amazon S3, Apache Airflow and SQL Server**.
- Implemented **data quality and observability** using **dbt tests orchestrated through Airflow**, catching upstream data defects before they reached reporting layers.
- Authored **AWS Glue** transformation scripts converting raw transactional data into curated analytics tables.
- Introduced **CI/CD practices** for pipeline deployment, reducing manual release effort and deployment errors.

Data Engineer -- Tredence Inc. GLOBAL HOSPITALITY GROUP

Bengaluru, India · Jul 2022 – Jan 2025 (2 yrs 7 mos)

- Delivered a full **migration from a legacy SAS platform to Snowflake** for **one of the world's largest hotel groups** -- 30+ brands across 130+ countries -- as part of the delivery team, owning several modules end to end.
- **Optimised legacy code and Snowflake queries**, improving pipeline runtime and reducing warehouse compute cost.
- Built **Python automation scripts** for account tracking and reconciliation, replacing recurring manual effort with a scheduled process.
- Developed **Snowflake stored procedures** and transformation logic serving downstream business users.
- Owned **daily client communication** for the engagement and ran **knowledge transfer and onboarding for junior engineers**.
- Recognised **twice with a Certificate of Appreciation** for exceptional delivery.

EDUCATION

Bachelor of Technology (B.Tech)

Birsa Institute of Technology (BIT), Sindri

CGPA 8.55 / 10

ADDITIONAL

Languages: English (Fluent) · Hindi (Native) · Urdu (Fluent) · Arabic (Reading & writing proficient; basic speaking and comprehension)

Notice Period: 15 days